

LVALUES AND RVALUES

- lvalues & rvalues aren't really language features.
- rather they are semantic properties of expressions.
- understanding them provides valuable insight into:
  - the behaviour of built-in operators.
  - the code generated to execute those operators.
  - the meaning of some otherwise cryptic compiler messages
  - reference types
  - overloaded operators

LVALUES AND RVALUES HAVE EVOLVED

- in early C, the concepts of lvalue & rvalue were fairly simple.
- early C++ added classes, const, & references.
- the concepts got more complicated.
- origins of lvalue & rvalue from historical persp.

## LVALUES

- in "the C programming language", Kernighan & ritchie wrote:
  - the name "lvalue" comes from the assignment expression

$E1 = E2$

in which left operand  $E1$  must be an lvalue expression.

- an lvalue is an expression referring to an object
- an object is a region of storage

## LVALUES AND RVALUES

int n; // a definition for an integer object  
named n

~~~

n=1; // an assignment expression

- n is a sub-expression referring to an int object.
- its an lvalue.
- i is a sub-expression not referring to an object.
- its an rvalue.
- an rvalue is simply an expression that's not an lvalue.

## LVALUES AND RVALUES

- $x[i+1] = \text{abs}(p \rightarrow \text{value});$
- $\underline{x[i+1]}$  is an expression.  
so is  $\text{abs}(p \rightarrow \text{value})$ .
- for the assignment to be valid:
  - the left operand must be an lvalue.  
it must refer to an object.
  - the right operand can either be an lvalue or rvalue.  
it can be any expression.

## A LOOK UNDER THE HOOD

- why make this distinction between lvalues & rvalues?
- so that compilers can assume that rvalues don't necessarily occupy storage.
- this offers considerable freedom in generating code for rvalue expressions.

int n; // a declaration for an integer object named n

~~~

n = 1; // an assignment expression

## DATA STORAGE FOR RVALUES

- a compiler might represent  $i$  as named data storage initialized with the value  $i$ , as if  $i$  were an lvalue.
- in assembly language, this might look like:

One : ; label for the following locat.  
word 1 ; allocate storage for holding  
the value 1.

- the compiler would generate code to copy from that initialized storage to the storage allocated for  $n$ .

mov n, one ; copy the value at one  
to location  $n$

## DATA STORAGE FOR RVALUES

- most machines provide instructions with an immediate operand:
- a source operand value can be part of an instruction.
- in assembly, this might look like:  
`mov n, #1 ; copy the value 1 to  
 location n`
- in this case:
  - the rvalue 1 never appears as an object in the data space.
  - rather, it appears as part of an instruction in the code space.

## DATA STORAGE FOR RVALUES

- On some machines, the preferred way to put the value 1 into an object might be to:
  - Clear the object
  - then increment it
- in assembly, this might look like:

```
clr n ; set n to zero
inc n ; increment n, effectively setting it
          to 1
```
- data representing the values 0 and 1 appear nowhere in either the source or object code.

MUST BE AN LVALUE == CANT BE AN  
RVALUE

`i = n;` // obviously silly

- this is trying to change the value of integer literal `i`.
- of course, C (and C++) reject it as error.
- but why exactly?
  - an assignment assigns a value to an object.
  - its left operand must be an lvalue.
  - but i is not an lvalue, its an rvalue.

## RECAP

- every expression in C++ is either an lvalue or an rvalue.
- in general:
  - lvalue is an expression that refers to an object.
  - rvalue is simply any expression that isn't an lvalue.
- caveat:

although this is true for non-class types in C++, it's not true for class types.

## LITERALS

- most literals are rvalues, including:
  - numeric literals, such as 3 & 3.14159
  - character literals, such as 'a'.
- they don't necessarily occupy data storage.
- however, string literals, such as "xyzzy", are lvalues.
- they occupy data storage.

## ENUMERATION CONSTANTS

- when used in expressions, enumeration constants are also rvalues.

```
enum color {red, green, blue};    c = green; ✓  
color c;  
                                blue = green; ✗
```

## LVALUES USED AS RVALUES

- an lvalue can appear on either side of assignment:

```
int m, n;
```

~~~

```
m = n; // OK: m & n are both lvalues
```

- this assignment uses the expression n as an lvalue.
- officially, C++ performs an lvalue-to-rvalue conversion.

## OPERANDS OF OTHER OPERATORS

- the concepts of lvalue and rvalue apply in all expressions.
  - not just assignment.
- for example, both the operands of the binary operator "+" must be expressions.
  - obviously, those expressions must have suitable types.
- but each operand can either be an lvalue or rvalue.

int x;

~~~

$\sim x + 2 \sim$ , // OK: lvalue + rvalue

$\sim 2 + x \sim$ ; // OK: rvalue + lvalue

## WHAT ABOUT THE RESULT?

- for built-in binary (non-assignment) operators such as "+":
  - the operands may be lvalues or rvalues.
  - but what about the result?
- an expression such as "m + n" places its result:
  - not in "m"
  - not in "n"
  - but rather in a compiler-generated temporary object, often a CPU register.
- Such temporary objects are rvalues.

## WHAT ABOUT THE RESULT?

- for example, this is (Obviously?) an error:

$m + 1 = n;$  // error... but why?

- the "+" operator has higher precedence than " $=$ ".
- thus, the assignment expression is equivalent to:

$(m + 1) = n;$  //error... but why?

- its an error because "m+i" yields an rvalue.

## UNARY &

- $\&e$  is a valid expression only if  $e$  is an lvalue.
  - thus,  $\&3$  is an error.
    - again, 3 does not refer to an object, so it is not addressable.
- although the operand must be an lvalue, the result is an rvalue.

int n, \*p;

~~~

$p = \&n;$  // OK: n is an lvalue

$\&n = p;$  // error:  $\&n$  is an rvalue

## UNARY \*

- in contrast to unary &, unary \* yields an lvalue.
- a pointer "p" can point to an object, so \*p is an lvalue.

```
int a[N];
```

```
int *p = a;
```

```
char *s = NULL; // = nullptr in modern C++
```

~ ~ ~

```
*p = 3; // OK
```

```
*s = '\0'; // undefined behaviour
```

- note: lvalue-ness is a compile-time property.
  - \*s is an lvalue even if "s" is null.
  - if "s" is null, evaluating \*s causes undefined behaviour.

## UNARY \*

- again, the result of the \* operator is an lvalue.
- however, its operand can be an rvalue.

$*(\text{p}+1) = 4; \quad // \text{OK}$

- here,  $\text{p}+1$  is an rvalue, but  $*(\text{p}+1)$  is an lvalue.
- the assignment stores the value 4 into the object referenced by  $*(\text{p}+1)$ .

## DATA STORAGE FOR EXPRESSIONS

- conceptually, rvalues (of non-class type) don't occupy data storage in the object program.
- in truth, some might.
- C & C++ insist that you program as if non-class rvalues don't occupy storage.
- conceptually, lvalues of any type occupy data storage.
  - in truth, the optimizer might eliminate some of them.
  - (but only when you won't notice).
- C & C++ let you assume that lvalues do occupy storage.

## NON-MODIFIABLE LVALUES

- in fact, not all lvalues can appear on the left of assignment.
- an lvalue is non-modifiable if it has a const-qualified type.

```
char const name[] = "dan";
```

~~~

```
name[0] = 'D'; //error name[0] is const
```

- `name[0]` is an lvalue, but its non-modifiable.
- each element of a `const` array is itself `const`.

## NON MODIFIABLE LVALUES

- lvalues & rvalues provide a vocabulary for describing subtle behavioral differences.
  - such as between enumerations constants & const objects.
- for example, this MAX is a constant of an unnamed enumeration type.

```
enum { MAX = 100 };
```
- unscoped enumeration values implicitly convert to integer.
- when MAX appears in an expression, it yields an integer rvalue.

```
MAX += 3; // error: MAX is an rvalue
```

```
int *p = &MAX; // error: again MAX is an rvalue
```

## NON MODIFIABLE LVALUES

- on the other hand, this MAX is a const-qualified object.

```
int const MAX = 100;
```

- when it appears in an expression, it's a non-modifiable lvalue.

- thus, you still can't assign to it.

```
MAX += 3; // error: MAX is non-modifiable
```

- however, you can take its address:

```
int const *p = &MAX; // OK: MAX is an lvalue
```

	can take address of	can assign to
lvalue	yes	yes
non-modifiable lvalue	yes	no
(non-class) rvalue	no	no

## CONST OBJECTS

- a const object is addressable.
- the compiler may generate storage to hold the const object's value.
- the compiler might find that the program never needs storage for a particular const object.
  - it often does.
  - in that case, the compiler need not allocate storage for that object.
- this behaviour for const objects is analogous to the behaviour for inline functions.

## REFERENCE TYPES

- the concepts of lvalues & rvalues help explain C++ reference types.
- references provide an alternative to pointers as a way of associating names with objects.
- C++ libraries often use references instead of pointers as function parameters and return types.

int i; // define i as an integer object

~~~

int &ri = i; // define ri as a "ref. to int"

- "ri" is an alias for "i."

## REFERENCE TYPES

- a reference is essentially a pointer that's automatically dereferenced each time it's used.
- you can rewrite most, if not all, code that uses a reference as a code that uses a const pointer, as in:

reference notation

int &ri = i;

ri = 4;

int j = ri + 2;

equivalent pointer notation

int \*const cpi = &i;

\*cpi = 4;

int j = \*cpi + 2;

- a reference acts like a const pointer that's dereferenced (has a \* in front of it) whenever you touch it.
- a reference yields an lvalue.

## INITIALIZING VS ASSIGNING

- initializing a reference associates the reference with an object.
- initializing a reference is also known as binding.
- assigning a reference stores through the reference and into the referenced object.

```
int &ri = i; // binds reference to object  
ri = 3;      // assigns to referenced object
```

## REFERENCES AND OVERLOADED OPERATORS

- what good are references?
- why not just use pointers.
- references can provide friendlier function interfaces.
- more specifically, C++ has references so that overloaded operators can look just like built-in operators...

## REFERENCES AND OVERLOADED OPERATORS

```
enum month {
```

```
    Jan, Feb, ..., Dec, month_end
```

```
};
```

```
typedef enum month month;
```

```
~~~
```

```
for (month m = Jan; m <= Dec; ++m) {
```

```
~~~
```

```
}
```

- this code compiles and executes as expected in C.
- however, it doesn't compile in C++...
- in C++, the built-in "`++`" won't accept an operand of enumeration type.
- you need to overload "`++`" for `month`.
- let's try it without references.

## REFERENCES AND OVERLOADED OPERATORS

```
void operator++(month &x) { // pass by value  
    x = static_cast<month>(x+1);  
}
```

- using this definition, `++m` compiles, but doesn't increment `m`.
  - it increments a copy of `m` in parameter `x`.
- also, this implementation lets you apply "`++`" to an rvalue, as in:  
`++Apr;` // compiles, but shouldn't
- a proper overloaded "`++`" should behave like the built-in "`++`", as in:  
`++42;` // compile error: can't increment an rvalue.

## REFERENCES AND OVERLOADED OPERATORS

- we need a "++" that passes in a month  
it can modify:

```
void operator ++(month *x) { // pass by  
    *x = static_cast<month>(*x + 1); // address  
}
```

- this works, but not like a builtin ++:

`++m;` // doesn't compile.

`++&m;` // compiles, increments m, but looks wrong

## REFERENCES AND OVERLOADED OPERATORS

- we really need a "++" that can modify a month object.
- but without passing explicitly by address

```
void operator++(month &x) { //pass by reference  
    x = static_cast<month>(x + 1);  
}
```

- using this definition:

```
++m; //compiles, increments m, and looks right
```

- as a bonus, this "+" operator won't accept an rvalue.

```
++Apr; //compile error
```

## REFERENCES AND OVERLOADED OPERATORS

- actually, a proper prefix "++" doesn't return void.
- it returns the incremented object by reference.

month & operator ++(month&n) { //non-void  
 return \*x = static\_cast<month>(x+1); return  
}

- this enables overloaded "++" to act even more like a built-in operator:

int j = ++m; //OK

month n = ++m; //OK

## "REFERENCE TO CONST" PARAMETERS

- just as you can have "pointer to const" parameters ..
- you can also have reference to const parameters.

`R f(T const &t);`

- a "reference to const" parameter will accept an argument that's either const or non-const.
- in contrast, a reference (to non-const) parameter will accept only non-const argument.
- when it appears in an expression, a "reference to const" yields ~~an~~ a non-modifiable lvalue.

## "REFERENCE TO CONST" PARAMETERS

- for the most part, a function declared as:

$R f(T \text{ const } \&t)$ ; // by "reference to const"

has the same outward behaviour as a function declared as:

$R f(T \&t)$ ; // by value

- that is, the calls look & act very much the same.

## "REFERENCE TO CONST" PARAMETERS

- either way, you declare "f", you write the argument expression the same way:

$T x;$

~~~

$f(x);$  // by value, or by "reference to const"?

- either way, calling "f" can't alter the actual argument, "x":

- by value: "f" has access only to a copy of "x", not "x" itself.
- by "reference to const": "f"'s parameter is declared to be non-modifiable.

## WHY USE "REFERENCE TO CONST"?

- why pass by "reference to const" instead of by value?
- passing by "reference to const" might be much more efficient than passing by value.
- it depends on the cost to make a copy.

## REFERENCES AND TEMPORARIES

- a "pointer to T" can only point to an lvalue of type T.
- similarly, a "reference to T" binds only to an lvalue of type T.
- for example, these are both compile errors:

int \*pi = &3; // can't apply & to 3

int &ri = 3; // can't bind this either

int i;

double \*pd = &i; // can't convert pointers

double &rd = i; // can't bind this either

## REFERENCES AND TEMPORARIES

- there's an exception to the rule that a reference must bind to an lvalue of referenced type:
  - a "reference to const T" can bind to an expression  $x$  that's not an lvalue of type T.
  - if there's a conversion from  $x$ 's type to T.
- in this case, the compiler creates a temporary object to hold a copy of  $x$  converted to T.
- this is so the reference has something to bind to.

## REFERENCES AND TEMPORARIES

- `int const & ri = 3;`
- When the program execution reaches this declaration, the program:
  1. creates an int temporary to hold the 3, and
  2. binds `ri` to the temporary.
- When execution leaves the scope containing `ri`, the program:
  3. destroys the temporary

## REFERENCES AND TEMPORARIES

double const &rd = ri;

- When the program execution reaches this declaration, the program:
  1. converts the value of ri from int to double.
  2. creates a temporary double to hold the converted result, and
  3. binds rd to the temporary.
- again, when execution leaves the scope containing rd, the program:
  4. destroys the temporary.

## REFERENCES AND TEMPORARIES

- this special behaviour enables passing by "reference to const" to consistently have the same outward behaviour as passing by value.

long double x;

void f(long double ld); // by value

~~~

f(x); // passes a copy of x

f(i); // passes a copy of i converted to long double

## REFERENCES AND TEMPORARIES

- this is the same example , except it uses a "reference to const" parameter in place of a value parameter.

long double x;

void f(long double const &ld); // by reference  
to const

~~~

f(x); // passes a reference to x

f(1); // passes a reference to a temporary  
containing 1 converted to long double.

- either way, the function calls behave the same.

## MIMICKING BUILT-IN OPERATORS

- recall the behaviour of the built-in + operator:
  - the operands may be lvalues or rvalues.
  - the result is always an rvalue.
- how do you declare an overloaded operator with the same behaviour?
- consider a rudimentary (character) string class with + as a concatenation operator.

## MIMICKING BUILT-IN OPERATORS

- you can declare operator + as a non-member:

```
class string {  
public:  
    string (String const &);  
    string (char const *); // converting constructor  
    string && operator= (String const &);  
    ~~~~  
};
```

string operator+ (String const &l, const String &r),

```
String s = "hello";  
String t = "world";  
~~~
```

lvalue or rvalue

s = s + ", " + t;

s = s + String (", ") + t; // lvalue + rvalue  
+ lvalue

Compiler applies converting constructor implicitly

## MIMICKING BUILT-IN OPERATORS

string operator+(string const &l, const string &r);

- the function returns its result by value.
- calling this operator + yields an rvalue.

string \*p = &(s+t); // error: can't take the address

## REFERENCES

- C++11 introduces another kind of reference.
- what C++03 calls "references", C++11 calls "lvalue references".
- this distinguishes them from C++11's new "rvalue references".
- except for the name change, lvalue references in C++11 behave just like references in C++03.

## RVALUE REFERENCES

- whereas an lvalue reference declaration uses the & operator, an rvalue reference uses the && operator.

- for example, this declares ri to be an "rvalue reference to int":

```
int &&ri = 10;
```

- you can use rvalue references as function parameters and return types, as in:

```
double &&f (int &&ri);
```

- you can also have an "rvalue reference to const", as in:

```
int const &&rci = 20;
```

## RVALUE REFERENCES

- rvalue references bind only to rvalues.
- this is true even for "rvalue reference to const".

```
int n = 10;
```

```
int && ri = n; //error: n is an lvalue
```

```
int const && rj = n; //error: n is an lvalue
```

## MOVE OPERATIONS

- modern C++ uses rvalue references to implement move operations that can avoid unnecessary copying:

```
class String S
```

```
public:
```

```
// copy operations
```

```
String (String const &); // constructor
```

```
String && operator= (String const &); // assignment
```

```
// move operations
```

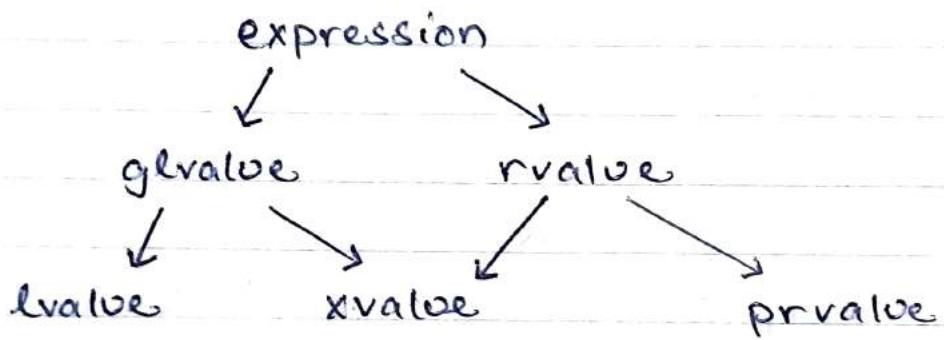
```
String (String &&) noexcept; // constructor
```

```
String && operator= (String &&) noexcept; // assignment
```

```
};
```

## VALUE CATEGORIES

- modern C++ introduces a more complex categorization of expressions:



- the newer categories are:

- glvalue: a "generalized" lvalue
- prvalue: a "pure" rvalue
- xvalue: an "expiring" lvalue

— X —